

AI Engagement Canvas

A practical discovery artifact for turning broad AI interest into a scopeable proof. Use in ideation, then carry the assumptions into estimate, build, evaluation, and retrospective.

1 · BUSINESS DECISION What decision, workflow, delay, quality failure, or recurring exception are we trying to improve? Who owns it today?	2 · AI LEVERAGE Should AI retrieve, predict, generate, classify, optimize, or act? What materially changes if it works?	3 · HUMAN AUTHORITY What must a person approve, review, override, escalate, or stop? What consequence of error is unacceptable?
4 · EVIDENCE & DATA What representative cases, context, labels, history, policies, and edge conditions are available? What is missing?	5 · PROOF DESIGN What is the smallest credible demonstration? Which eval cases, failure modes, and business measures will decide whether to continue?	6 · PRODUCTION HANDOFF What integrations, permissions, observability, rollback, support, documentation, training, and owner cadence are required?

EXPECTED VALUE Economic or mission importance	DATA READINESS Representative context available	ADOPTION BURDEN Behavior and workflow change	RISK / CONSEQUENCE Failure cost and control needs	TIME TO EVIDENCE Speed to a decision-quality proof
---	---	--	---	--

Decision gate: Proceed, narrow, pause for data/controls, or stop. This canvas is a candidate operating hypothesis, not a representation of Capitalize's proprietary methodology.

