

Make AI credible in the room—and useful in the workflow.

A full-lifecycle operating thesis for the AI Senior Consultant role at Capitalize Analytics.

Mandate, company moment, and operating context

01 /
03

Actual mandate

Turn ambiguous executive AI demand into well-scoped, technically credible, governed work that produces evidence, survives production handoff, and strengthens a growing consulting practice. The posting explicitly expects one person to lead ideation, advise executives, support sales, scope engagements, build solutions, and present results.

Why now

DEFINED ENTRY

A 90-minute ideation session creates a repeatable front door for clients who need direction.

TIMED PROOF

A 45-day AI-to-ROI model makes scope, assumptions, governance, and evidence commercially visible.

PLATFORM DEPTH

Claude, Databricks, Snowflake, cloud platforms, LLMs, and agentic frameworks widen both possibility and consequence.

PRACTICE LEVERAGE

The role must connect market promise, delivery quality, reusable knowledge, and client ownership.

Operating tensions

AMBITION / EVIDENCE

Executives want speed; delivery needs proof that survives real data, users, and controls.

SPEED / GOVERNANCE

Fast evidence matters, but human authority, permission, and escalation cannot be bolted on later.

BESPOKE / REUSABLE

Client specificity creates value; disciplined learning creates practice leverage.

EXPERTISE / OWNERSHIP

The consultant must add capability without creating permanent dependence.

Stakeholder topology

Executive sponsors and business owners · AI & Advanced Analytics practitioners · data engineering and platform teams · automation and BI practitioners · account growth and presales · strategic alliances and partners · security, governance, legal, and compliance · client frontline users and operating owners.

Decisions the role should improve

- Which problem deserves an AI experiment, and what is the smallest proof that can change the next decision?
- What assumptions must be known before estimating, and what remains a human decision?
- Which controls are mandatory before production, what becomes reusable practice IP, and when should the team stop?

The AI Proof Ledger and evidence map

02 / 03

Thesis

Strategy and delivery should form one continuous proof chain. The Senior Consultant creates leverage by preserving decision context, human authority, evidence standards, handoff obligations, and client ownership across the full engagement.

Six continuity records

01 · DECISION RECORD

Outcome, workflow, consequence, and accountable owner.

02 · LEVERAGE HYPOTHESIS

Expected change in economics, speed, quality, or risk—and what could disprove it.

03 · AUTHORITY BOUNDARY

Human decisions, reviews, overrides, escalations, and stop conditions.

04 · EVIDENCE STANDARD

Representative cases, quality thresholds, business measures, and failure modes.

05 · HANDOFF CONTRACT

Integrations, permissions, observability, support, rollback, and training.

06 · OWNERSHIP CADENCE

Who reviews performance, exceptions, risk, and improvement after delivery.

Verified evidence map

VAPE-JET

Direct AI-first operating transformation across production, quality, support, customer success, ERP, engineering issue flow, telemetry, and AI-assisted workflows.

DUDEWORTH

AI-readiness and agentic workflow frameworks grounded in value, feasibility, governance, evaluation, and adoption burden.

AMAZON

Launch, execution, standard work, escalation, and customer-backward operations in a 24/7 high-velocity environment.

COMPUNETICS

High-reliability engineering, manufacturing, quality, customer delivery, process controls, and technical communication.

Balanced scorecard

- **Commercial:** ideation conversion, proposal cycle, estimate accuracy, expansion path.
- **Delivery:** time to evidence, rework, evaluation pass rate, handoff quality.
- **Client value:** cycle time, quality, cost, risk, throughput, or revenue tied to the use case.
- **Trust and adoption:** groundedness, overrides, escalation, target-workflow use, training, and owner cadence.
- **Practice leverage:** reusable discovery assets, eval sets, reference patterns, demos, and enablement.

Objection, entry plan, questions, and sources

03 /
03

Strongest hiring objection

Concern: Is Russell current and hands-on enough for a role that expects production-ready AI delivery across modern data platforms—not only strategy and operations leadership?

Honest response: Do not claim production ownership of every platform. Verified toolkit: Python, SQL, TensorFlow, LLM/RAG patterns, agentic workflows, human-in-the-loop design, and direct AI-assisted operating transformation. Test the bridge with a practical case, architecture review, and pair-build exercise.

First 90 days

DAYS 1–30 · READ

Shadow ideation, demos, scoping, delivery, handoff, and retrospectives. Map decision rights and repeated friction.

DAYS 31–60 · CO-DELIVER

Turn one opportunity into a decision brief, evaluation set, governed prototype, and handoff checklist.

DAYS 61–90 · LEAD

Lead a bounded workstream, run the retrospective, and establish client ownership cadence.

COMPOUND

Convert validated learning into one reusable asset, one enablement session, and one stop condition.

Executive questions

1. Where does the practice lose the most momentum between ideation, scope, proof, handoff, and adoption?
2. What outcome matters most in the first six months: delivery, presales conversion, reusable IP, partner motion, or a balanced mix?
3. How are decision rights divided when an opportunity crosses AI, data engineering, automation, account growth, and alliances?
4. Which parts of the 45-day model are standardized, and where should a Senior Consultant exercise judgment?
5. What would make leadership say this person made the practice easier to sell, deliver, and trust?

Sources and evidence boundary

Job posting: <https://ats.rippling.com/capitalize-data-analytics-llc/jobs/014035f0-0e5f-47f8-970c-53906fa7d152>

Practice: <https://capitalizeconsulting.com/ai-and-advanced-analytics/>

Company: <https://capitalizeconsulting.com/who-we-are/>

Candidate claims come from RoleForge's verified evidence. Company-specific recommendations remain hypotheses for discovery unless supported by the public sources above.