

Make AI credible in the room—and useful in the workflow.

Mandate in one sentence: Turn ambiguous executive AI demand into well-scoped, technically credible, governed work that produces evidence, survives handoff, and strengthens a growing consulting practice.

Why this role exists now

Capitalize's AI & Advanced Analytics practice is growing and the posting explicitly spans ideation, executive advisory, sales demos, scoping, estimation, solution design, building, POCs, and production-ready deliverables. The company's AI service line already offers a 90-minute ideation session and a 45-day proof-to-working-model engagement. In June 2026, Capitalize announced a dedicated Anthropic practice and Select Tier status in the Claude Partner Network Services Track, alongside production Claude deployments.

My inference about the real mandate

The Senior Consultant is not only an individual contributor on AI delivery. The role is a connective tissue position: reduce friction among client discovery, presales credibility, scope quality, technical proof, production handoff, adoption, and reusable practice learning.

White-space thesis

Strategy and delivery should form one continuous proof chain. Begin with the business decision, identify the leverage point, define human authority, build the narrowest credible proof, instrument it, engineer the handoff, and leave the client with a capability it can own.

Visual metaphor

The campaign uses an "AI Proof Ledger"—one live continuity record that preserves the business decision, leverage hypothesis, human authority, evidence standard, handoff obligations, and client ownership cadence across the engagement. It is intentionally structured as a ledger rather than a proving ground, bench, graph, flywheel, constellation, digital spine, or control room.

Operating tensions

Tension	Why it matters	Operating response
Executive ambition vs delivery evidence	AI demand often arrives as broad aspiration; the role must create a narrow, testable decision without killing momentum.	Discovery anchored in a business decision, evidence source, owner, consequence of error, and measurable proof criterion.
Speed vs governance	The 45-day offer creates useful urgency, while production credibility requires security, evaluation, authority boundaries, and traceability.	Governance designed into the proof: permissions, eval cases, escalation, logging, stop conditions, and named human authority.
Bespoke consulting vs repeatable practice leverage	Clients need context-specific solutions, but a growing practice cannot re-invent discovery, evaluation, and handoff every time.	Standardize decision artifacts and quality gates; customize the business logic, data, integration, and adoption model.
Presales promise vs delivery reality	The role presents demos and scopes work, then stays close to the build. That creates an opportunity to improve estimate truthfulness.	Use proof plans and assumptions that remain visible through build and retrospective.
Tool breadth vs problem discipline	Capitalize works across Claude, OpenAI, Databricks, Snowflake, Fabric, automation, and BI ecosystems.	Choose technology after the workflow, decision, constraints, and evidence needs are explicit.

Decision architecture

- Which client problem deserves an AI experiment?
- What is the smallest proof that can change a funding or scope decision?
- What must be known before giving an estimate?
- Which controls are mandatory before external users or production tools are involved?
- What remains a human decision?
- What learning becomes reusable practice IP—and what must remain client-specific?
- When should the team stop, narrow, or escalate rather than continue building?

The AI Proof Ledger

Station	Core question	Working artifact	Failure avoided
Decision record	What business decision, delay, quality failure, or exception is worth changing, and who owns it?	Decision and workflow record	Solution looking for a problem
Leverage hypothesis	Where can AI materially change economics, speed, quality, or risk?	Leverage statement plus decision threshold	Interesting demo with no business consequence
Authority boundary	What must a person decide, review, override, escalate, or stop?	Authority and escalation record	Automation that obscures accountability
Evidence standard	What is the narrowest credible proof?	Representative cases, evaluation rubric, failure-mode log	Demo success mistaken for production readiness
Handoff contract	What must exist for safe operation and adoption?	Integration, permission, observability, support, training, rollback, and ownership checklist	Orphaned proof or uncontrolled automation
Ownership cadence	How does the client own and improve the result?	Decision records, documentation, enablement, and review cadence	Permanent dependency with weak adoption

Balanced scorecard

Dimension	Measures to baseline with leadership
Commercial	Ideation-to-engagement conversion; expansion path; estimate accuracy; proposal cycle time.
Delivery	Time to first working evidence; rework; evaluation pass rate; production handoff quality.
Client value	Use-case-specific cycle time, quality, cost, risk, throughput, or revenue measure.
Trust	Groundedness, override patterns, escalation quality, permission failures, auditability, incidents.
Adoption	Use in target workflow, training, abandonment points, confidence, ownership cadence.
Practice leverage	Reusable discovery assets, evaluation sets, reference patterns, demos, practitioner enablement.

Verified evidence → role relevance

Evidence	Verified fact pattern	Transfer to Capitalize
Vape-Jet	Direct AI-first operating transformation across production, quality, customer/support, engineering issue flow, ERP/Odoo, Jira, HubSpot, Slack, telemetry, and AI-assisted workflows.	Understands the operating middle between AI possibility and adopted workflow change.
DudeWorth	Direct work on AI-readiness frameworks and agentic workflow patterns using context, retrieval, human judgment, escalation, decision records, evaluation, and governance rails.	Structured discovery, use-case prioritization, governed agent design, and executive translation.
Amazon	Led launch and execution in a 24/7 customer-backward environment supporting approximately five million second-day orders annually; built operating routines, visibility, escalation, and standard work.	Mechanism discipline, measurable operations, urgency without loss of accountability.
Compunetics	Led engineering, manufacturing, quality, customer delivery, and complex technical programs; built controls and root-cause practices; worked with high-consequence partners.	Technical credibility, consequence-aware delivery, and communication across executive, engineering, and operations contexts.
ZeusVu	Led regulated autonomous drone operations with a 100% safety record and no FAA incidents; developed TensorFlow-based concepts for aerial inventory and computer vision.	Autonomy, risk boundaries, field data, and honest distinction between concepts and deployments.

Strongest objection

Concern: Is Russell current and hands-on enough for a senior consulting role that expects production-ready AI delivery across modern platforms?

Response: Do not overclaim. The verified toolkit includes Python, SQL, TensorFlow, LLM/RAG patterns, agentic workflows, HITL design, and direct AI-first operating transformation. The differentiator is technical fluency combined with accountability for real workflows, controls, adoption, and consequences. Offer a practical case, architecture review, or pair-build exercise as the proof mechanism.

Entry plan

Days 1–30 · Read the system

- Shadow ideation, sales demos, scoping, delivery, and retrospectives.
- Map lead signal → use-case choice → estimate → proof → evaluation → handoff → adoption → expansion.
- Learn reference architectures, security/governance patterns, partner motions, and escalation paths.
- Build a practice friction register and decision-rights map.

Days 31–60 · Co-deliver the proof

- Co-lead ideation and convert one opportunity into a decision brief with value, data, risk, human authority, and proof criteria.
- Build or pair-build a reusable demo/reference pattern within the sanctioned stack.
- Test a scoping and production-handoff checklist.
- Instrument value and adoption measures at engagement start.

Days 61–90 · Lead and teach

- Lead a bounded engagement or major workstream end to end with leadership support.
- Retrospect against estimate, proof criteria, client adoption, and reusable practice learning.
- Turn validated learning into one enablement session, one reusable asset, and one documented stop condition.
- Propose the next practice improvement only after evidence reveals leverage.

Change burden: The thesis works only if sellers, advisors, builders, and client owners keep assumptions visible across handoffs; if evaluation and adoption are defined before the demo; and if client teams accept explicit ownership of the workflow after delivery.

Executive questions

1. Where does the AI practice lose the most momentum between ideation, scope, proof, production handoff, and adoption?
 2. What outcome matters most for this hire in the first six months: delivery, presales conversion, reusable IP, partner motion, or a deliberate mix?
 3. How are decision rights divided across AI, Data Engineering, Process Automation, account growth, and strategic alliances when opportunities cross practices?
 4. What have the first production Claude deployments taught the team that should change ideation or scoping?
 5. Which parts of the 45-day model are intentionally standardized, and where should a Senior Consultant exercise judgment?
 6. When a proof is technically successful but behavior does not change, who owns the adoption problem?
 7. What would make leadership say after one year that this person made the practice easier to sell, deliver, and trust?
-

Primary research sources

- Capitalize Analytics AI Senior Consultant posting (Rippling)
- Capitalize Analytics — AI & Advanced Analytics service page
- Capitalize Analytics — Claude Partner Network announcement, June 11, 2026
- Capitalize Analytics — Who We Are
- Capitalize Analytics current careers board

Research integrity note: Company-specific recommendations in this brief are hypotheses for discovery unless explicitly stated as public facts. Candidate claims are restricted to verified RoleForge evidence.